

End-to-End Kubernetes and AWS Infrastructure Deployment with Terraform: A Static Website on NGINX

Project Overview

This project demonstrates the deployment of a static website using AWS infrastructure provisioned via Terraform, containerized using Docker, and deployed on a Kubernetes cluster using Minikube. The NGINX web server is used to serve the static website, and the setup is exposed using Kubernetes Ingress for external access.

Goals

- Infrastructure:

Provision an AWS VPC, Subnets, EC2 instances, and an Elastic Load Balancer (ELB) using Terraform.

- Application:

Create a static website, serve it using NGINX, and dockerize the application.

- Deployment:

Deploy the containerized application using Kubernetes (Minikube) with 4 replicas, exposing the service using ClusterIP and Ingress.

Technologies Used

- Terraform: For Infrastructure as Code (IaC) on AWS.
- AWS: VPC, Subnets, EC2, ELB.
- Docker: Containerization of the static website.
- Kubernetes (Minikube): Container orchestration for managing the application.
- NGINX: Web server for serving the static content.

Infrastructure Setup with Terraform

Description:

We will provision an AWS Virtual Private Cloud (VPC) with two subnets spread across two Availability Zones. Additionally, EC2 instances and an Elastic Load Balancer (ELB) will be created to distribute traffic.

Steps:

1. Install Terraform: if you haven't already.
 - Terraform Version: 1.4+
 - AWS Provider Plugin
2. Terraform Code: Create a main.tf file for the following setup:
 - VPC Creation
 - Subnet Allocation (2 AZs)
 - EC2 Instances and Security Groups
 - Elastic Load Balancer (ELB) setup
 - Ensure to manage Terraform state locally using terraform.tfstate file.

Terraform Configuration (Sample):

```
provider "aws" {  
    region = "us-east-1"  
    access_key = "*****"  
    secret_key = "*****"  
}  
  
# VPC  
resource "aws_vpc" "main" {
```

```

    cidr_block = "10.0.0.0/16"
}

#subnets
resource "aws_subnet" "subnet_a" {
    vpc_id = aws_vpc.main.id
    cidr_block = "10.0.1.0/24"
    availability_zone = "us-east-1a"
}

resource "aws_subnet" "subnet_b" {
    vpc_id = aws_vpc.main.id
    cidr_block = "10.0.2.0/24"
    availability_zone = "us-east-1b"
}

#internet Gateway

resource "aws_internet_gateway" "igw" {
    vpc_id = aws_vpc.main.id
}

# Route table

resource "aws_route_table" "main_rt" {
    vpc_id= aws_vpc.main.id
    route {
        cidr_block= "0.0.0.0/0"
        gateway_id = aws_internet_gateway.igw.id
    }
}

# secutiry group
resource "aws_security_group" "web_sg" {
    vpc_id = aws_vpc.main.id
    ingress {
        from_port = 80
        to_port = 80
        protocol = "tcp"
        cidr_blocks = ["0.0.0.0/0"]
    }
}

```

```

    }
    ingress {
      from_port = 22
      to_port = 22
      protocol = "tcp"
      cidr_blocks = ["0.0.0.0/0"]
    }
  }
}

# ELB
resource "aws_lb" "app_lb" {
  name = "app-load-balancer"
  internal = false
  load_balancer_type = "application"
  subnets = [aws_subnet.subnet_a.id, aws_subnet.subnet_b.id]
  security_groups = [aws_security_group.web_sg.id]
}

```

Commands:

- Initialize and apply Terraform configuration:

terraform init

terraform apply

```

Windows PowerShell
PS C:\Users\Fr3on\Desktop\Project-1> terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.67.0

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
PS C:\Users\Fr3on\Desktop\Project-1>

```

```
PS C:\Users\Fr3on\Desktop\Project-1> terraform apply

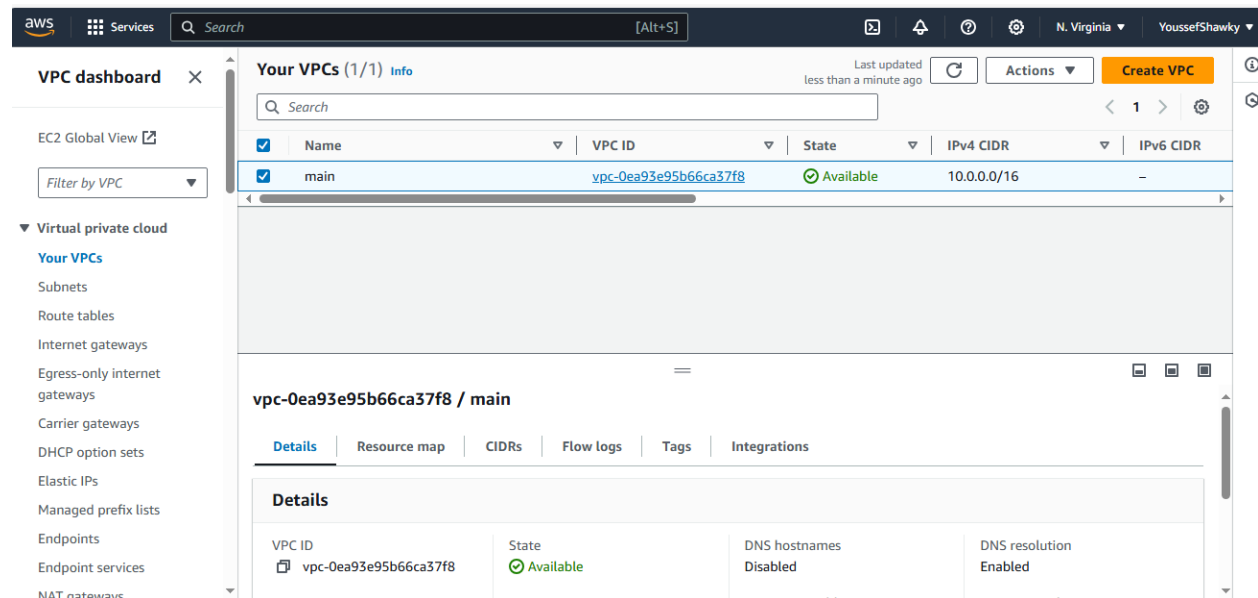
Terraform used the selected providers to generate the following execution plan. Resource actions are
indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_internet_gateway.igw will be created
+ resource "aws_internet_gateway" "igw" {
+   arn          = (known after apply)
+   id           = (known after apply)
+   owner_id     = (known after apply)
+   tags_all     = (known after apply)
+   vpc_id       = (known after apply)
}

# aws_lb.app_lb will be created
+ resource "aws_lb" "app_lb" {
+   arn                  = (known after apply)
+   arn_suffix           = (known after apply)
+   client_keep_alive    = 3600
+   desync_mitigation_mode = "defensive"
+   dns_name              = (known after apply)
+   drop_invalid_header_fields = false
+   enable_deletion_protection = false
}
```

AWS console verification:



aws

Services

Search

[Alt+S]

N. Virginia

YoussefShawky

VPC dashboard

EC2 Global View

Filter by VPC

Virtual private cloud

Your VPCs

Subnets

Route tables

Internet gateways

Egress-only internet gateways

Carrier gateways

DHCP option sets

Elastic IPs

Managed prefix lists

Endpoints

Endpoint services

NAT gateways

Subnets (2/2) Info

Last updated 4 minutes ago

Actions

Create subnet

Find resources by attribute or tag

<input checked="" type="checkbox"/>	Name	Subnet ID	State	VPC	IP
<input checked="" type="checkbox"/>	Subnet_a	subnet-049387a9ca8f210d3	Available	vpc-0ea93e95b66ca37f8 main	10
<input checked="" type="checkbox"/>	Subnet_b	subnet-0b4f2b996dee69b24	Available	vpc-0ea93e95b66ca37f8 main	10

Subnets: subnet-049387a9ca8f210d3, subnet-0b4f2b996dee69b24

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

aws

Services

Search

[Alt+S]

N. Virginia

YoussefShawky

VPC dashboard

EC2 Global View

Filter by VPC

Virtual private cloud

Your VPCs

Subnets

Route tables

Internet gateways

Egress-only internet gateways

Carrier gateways

DHCP option sets

Elastic IPs

Managed prefix lists

Endpoints

Endpoint services

NAT gateways

Internet gateways (1) Info

Actions

Create internet gateway

Search

Name	Internet gateway ID	State	VPC ID
-	igw-08198fe9f45b2b03c	Attached	vpc-0ea93e95b66ca37f8 main

Select an internet gateway above

Services

Search

[Alt+S]

N. Virginia

YoussefShawky

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Trust Stores

Auto Scaling

EC2 > Load balancers

Load balancers (1/1)

Elastic Load Balancing scales your load balancer capacity automatically in response to changes in incoming traffic.

Filter load balancers

< 1 >

<input checked="" type="checkbox"/>	Name	DNS name	State	VPC ID	Availability Zones
<input checked="" type="checkbox"/>	app-load-balancer	app-load-balancer-14645...	Active	vpc-0ea93e95b66ca37f8	2 Availability Zones

Load balancer: app-load-balancer

Details

Load balancer type	Status	VPC	Load balancer IP address type
Application	Active	vpc-0ea93e95b66ca37f8	IPv4
Scheme	Hosted zone	Availability Zones	Date created
Internet-facing	Z35SXDOTRQ7X7K	subnet-0b4f2b996dee69b24 us-east-1a (use1-az1)	September 20, 2024, 16:45 (UTC+03:00)

CloudShell

Feedback

© 2024, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

Building the Docker Image (Static Website + NGINX)

Description:

This step involves creating a simple static website (with home, info, and content pages), serving it with NGINX, and building a Docker image.

Steps:

1. Create a Static Website:

- Files:
 - index.html (Home page)
 - info.html (Info page)
 - content.html (Content page)

2. Create a Dockerfile for NGINX:

- This will copy your static website into NGINX's web root.

Dockerfile:

```
FROM nginx:alpine
COPY ./static-site /usr/share/nginx/html
EXPOSE 80
```

3. Build and Push the Docker Image:

```
docker build -t youssefshawky/project-1 .
```

```
docker push youssefshawky/project-1
```


New More Docker. Easy Access. New Streamlined Plans. Learn more. →

dockerhub

Explore

Repositories

Organizations

Usage

ctrl+K

Y

youssefshawky / Repositories / project-1 / General

Using 0 of 1 private repositories.

General

Tags

Builds

Collaborators

Webhooks

Settings

youssefshawky/project-1

Last pushed 7 days ago

This repository does not have a description

This repository does not have a category

Docker commands

To push a new tag to this repository:

docker push youssefshawky/project-1:tagname

Public View

Tags

This repository contains 1 tag(s).

Tag	OS	Type	Pulled	Pushed
latest		Image	5 days ago	7 days ago

Automated Builds

Manually pushing images to Hub? Connect your account to GitHub or Bitbucket to automatically build and tag new images whenever your code is updated, so you can focus your time on creating.

Available with Pro, Team and Business subscriptions. [Read more about automated builds](#)

Kubernetes Setup (Minikube)

Description:

Deploy the static website Docker image on a Kubernetes cluster using Minikube. Create a deployment with 4 replicas, expose it using ClusterIP, and use Ingress to expose it externally.

Steps:

```
# Create a deployment with 4 replicas

apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-website
spec:
  replicas: 4
  selector:
    matchLabels:
      app: my-website
  template:
    metadata:
      labels:
        app: my-website
    spec:
      containers:
        - name: my-website
          image: youssefshawky/project-1:latest
          ports:
            - containerPort: 80

---

# Exposing the deployment

apiVersion: v1
kind: Service
metadata:
  name: my-website
spec:
  selector:
    app: my-website
  ports:
```

```

- protocol: TCP
  port: 80
  targetPort: 80
  type: ClusterIP

---
# Set the ingress and expose the front end service to mapped to the backend
service

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-website-ingress
spec:
  rules:
    - host: mywebsite.io
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: my-website
                port:
                  number: 80

```

Minikube kubectl command verification:

Windows PowerShell

```

PS C:\Users\Fr3on\Desktop\Project-1> kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
my-website-6cffd4ff6c-df26s        1/1     Running   3 (5d ago)  6d23h
my-website-6cffd4ff6c-n27b6        1/1     Running   3 (5d ago)  6d23h
my-website-6cffd4ff6c-nbxj4        1/1     Running   3 (5d ago)  6d23h
my-website-6cffd4ff6c-t6svf        1/1     Running   3 (5d ago)  6d23h
PS C:\Users\Fr3on\Desktop\Project-1>

```


Windows PowerShell

```

PS C:\Users\Fr3on\Desktop\Project-1> kubectl get deploy
NAME            READY   UP-TO-DATE   AVAILABLE   AGE
my-website      4/4     4             4           6d23h
PS C:\Users\Fr3on\Desktop\Project-1>

```

```
PS C:\Users\Fr3on\Desktop\Project-1> kubectl get svc
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes          ClusterIP     10.96.0.1     <none>         443/TCP    6d23h
my-website          ClusterIP     10.97.185.41  <none>         80/TCP     6d23h
PS C:\Users\Fr3on\Desktop\Project-1>
```

 Windows PowerShell

```
PS C:\Users\Fr3on\Desktop\Project-1> kubectl get ingress
NAME                CLASS    HOSTS          ADDRESS          PORTS    AGE
my-website-ingress  nginx   mywebsite.io   192.168.49.2    80       6d23h
PS C:\Users\Fr3on\Desktop\Project-1>
```

Troubleshooting Section

Issues Faced:

- Ingress Not Working Externally:

Despite applying the Ingress configuration, the website wasn't accessible from the browser on the local machine. When running curl inside the Minikube VM (via SSH), the service was reachable, but not externally.

```
docker@minikube: ~  
PS C:\Users\Fr3on\Desktop\Project-1> minikube ssh  
W0920 17:21:51.271559 3460 main.go:291] Unable to resolve the current Docker CLI context "default": context "default":  
context not found: open C:\Users\Fr3on\.docker\contexts\meta\37a8eec1ce19687d132fe29051dca629d164e2c4958ba141d5f4133a33f06  
88f\meta.json: The system cannot find the path specified.  
docker@minikube:~$ curl mywebsite.io  
<!DOCTYPE HTML>  
<!--  
Forty by HTML5 UP  
html5up.net | @ajlkn  
Free for personal and commercial use under the CCA 3.0 license (html5up.net/license)  
-->  
<html>  
  <head>  
    <title>Forty by HTML5 UP</title>  
    <meta charset="utf-8" />  
    <meta name="viewport" content="width=device-width, initial-scale=1, user-scalable=no" />  
    <link rel="stylesheet" href="assets/css/main.css" />  
    <noscript><link rel="stylesheet" href="assets/css/noscript.css" /></noscript>  
  </head>  
  <body class="is-preload">  
  
    <!-- Wrapper -->  
    <div id="wrapper">  
  
      <!-- Header -->  
      <header id="header" class="alt">  
        <a href="index.html" class="logo"><strong></strong> <span></span></a>  
        <nav>  
          <a href="#menu">Menu</a>
```

What I've Tried So Far:

- Verified the Ingress controller and services are running as expected.
- Enabled the Minikube Ingress addon.
- Tried using port-forwarding and ``minikube tunnel``, but the Ingress remains inaccessible externally.
- Configured custom routes to the Minikube subnet, but no luck.

Interestingly, pinging from Minikube to my local machine works, but not the other way around.

My Thoughts:

Minikube seems to run as a "closed cluster," possibly by design. I suspect this networking isolation is why I can't access the Ingress externally.

Future Work

CI/CD Pipelines:

- Jenkins Setup:

Future improvements include automating the build and deployment process using Jenkins. This would allow continuous integration of code changes and their deployment to the Kubernetes cluster.